

Detection of Non-functional Water Wells Using Machine Learning Algorithms



Overview

Tanzania, as a developing country, struggles with providing clean water to its population of over 57,000,000. There are many wells (water points) already established in the country, but some are in need of repair while others have failed altogether. The Government of Tanzania is looking to find patterns in non-functional wells to influence how new wells are built.

Business Problem

Tanzania faces challenges in ensuring access to clean water for its population due to non-functional water wells. We will be looking to predict patterns in non-functional wells to inform more robust construction methods for new wells.

Data Understanding

Dataset

Driven Data - Tanzanian Water Wells

- [Labels](#)
- [Values](#)

Target

- functional : the well is operational and there are no repairs needed
- functional needs repair : the well is operational, but needs repairs
- non functional : the well is not operational

Features

- amount_tsh : Total static head (amount water available to well)
- date_recorded : The date the row was entered
- funder : Who funded the well
- gps_height : Altitude of the well
- installer : Organization that installed the well
- longitude : GPS coordinate
- latitude : GPS coordinate
- wpt_name : Name of the well if there is one

- num_private :Private use or not
- basin : Geographic water basin
- subvillage : Geographic location
- region : Geographic location
- region_code : Geographic location (coded)
- district_code : Geographic location (coded)
- lga : Geographic location
- ward : Geographic location
- population : Population around the well
- public_meeting : True/False
- recorded_by : Group entering this row of data
- scheme_management : Who operates the well
- scheme_name : Who operates the well
- permit : If the well is permitted
- construction_year : Year the well was constructed
- extraction_type : The kind of extraction the well uses
- extraction_type_group : The kind of extraction the well uses
- extraction_type_class : The kind of extraction the well uses
- management : How the well is managed
- management_group : How the well is managed
- payment : What the water costs
- payment_type : What the water costs
- water_quality : The quality of the water
- quality_group : The quality of the water
- quantity : The quantity of water
- quantity_group : The quantity of water
- source : The source of the water
- source_type : The source of the water
- source_class : The source of the water
- waterpoint_type : The kind of well
- waterpoint_type_group : The kind of well

Limitations

- Do not know what types of repairs are needed
- Not able to use the well age, due to missing construction years
- Without full population information, we do not know the supply needs

Data Preparation

Import and Read Datasets

```
In [3]: # Import standard packages
import pandas as pd
import numpy as np
```

```

import matplotlib.pyplot as plt
import seaborn as sns
%matplotlib inline
import datetime

# Model Selection
from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score

# Classification Models
from sklearn.linear_model import LogisticRegression
from sklearn.tree import DecisionTreeClassifier
from sklearn.dummy import DummyClassifier
from sklearn.ensemble import RandomForestClassifier

from sklearn.metrics import plot_confusion_matrix, accuracy_score, f1_score, precision_score
from sklearn.metrics import classification_report, confusion_matrix
from sklearn.metrics import roc_curve, auc, roc_auc_score

# Scalers
from sklearn.impute import SimpleImputer
from sklearn.preprocessing import StandardScaler

# Categorical Create Dummies
from sklearn.preprocessing import OneHotEncoder

# Column Transformer
from sklearn.compose import ColumnTransformer

# Pipeline
from sklearn.pipeline import Pipeline

# Base
from sklearn.base import BaseEstimator, TransformerMixin

```

```

In [4]: # Load the datasets
features_data = 'data/WellWaterData.csv'
target_data = 'data/TargetData.csv'

features = pd.read_csv(features_data)
target = pd.read_csv(target_data)

# Display the contents of the datasets
features.head(), target.head()

```

```

Out[4]: (
   id  amount_tsh  date_recorded  funder  gps_height  installer \
0  69572      6000.0   2011-03-14   Roman      1390      Roman
1   8776         0.0   2013-03-06  Grumeti      1399    GRUMETI
2  34310        25.0   2013-02-25 Lottery Club      686  World vision
3  67743         0.0   2013-01-28   Unicef      263     UNICEF
4  19728         0.0   2011-07-13 Action In A         0     Artisan

   longitude  latitude  wpt_name  num_private  ...  payment_type \
0  34.938093  -9.856322      none           0  ...    annually
1  34.698766  -2.147466   Zahanati           0  ...    never pay
2  37.460664  -3.821329  Kwa Mahundi           0  ...  per bucket
3  38.486161 -11.155298  Zahanati Ya Nanyumbu           0  ...    never pay
4  31.130847  -1.825359    Shuleni           0  ...    never pay

   water_quality  quality_group  quantity  quantity_group \
0             soft           good      enough           enough
1             soft           good  insufficient  insufficient

```

	2	3	4
	soft	good	enough
	soft	good	dry
	soft	good	seasonal

	source	source_type	source_class
0	spring	spring	groundwater
1	rainwater harvesting	rainwater harvesting	surface
2	dam	dam	surface
3	machine dbh	borehole	groundwater
4	rainwater harvesting	rainwater harvesting	surface

	waterpoint_type	waterpoint_type_group
0	communal standpipe	communal standpipe
1	communal standpipe	communal standpipe
2	communal standpipe multiple	communal standpipe
3	communal standpipe multiple	communal standpipe
4	communal standpipe	communal standpipe


```
[5 rows x 40 columns],
   id      status_group
0  69572      functional
1   8776      functional
2  34310      functional
3  67743  non functional
4  19728      functional)
```

Merging Datasets

```
In [5]: # Checking for unique IDs in both datasets to ensure they match
unique_ids_features = features['id'].nunique()
unique_ids_target = target['id'].nunique()

unique_ids_features, unique_ids_target
```

```
Out[5]: (59400, 59400)
```

```
In [6]: # Merging the datasets on the 'id' column
merged_data = pd.merge(features, target, on='id')

# Displaying the first few rows of the merged dataset
merged_data.head()
```

Out [6]:

	id	amount_tsh	date_recorded	funder	gps_height	installer	longitude	latitude	wpt
0	69572	6000.0	2011-03-14	Roman	1390	Roman	34.938093	-9.856322	
1	8776	0.0	2013-03-06	Grumeti	1399	GRUMETI	34.698766	-2.147466	Z
2	34310	25.0	2013-02-25	Lottery Club	686	World vision	37.460664	-3.821329	N
3	67743	0.0	2013-01-28	Unicef	263	UNICEF	38.486161	-11.155298	Z
4	19728	0.0	2011-07-13	Action In A	0	Artisan	31.130847	-1.825359	Nar

5 rows x 41 columns

In [7]: merged_data.info()

```

<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 41 columns):
#   Column                                Non-Null Count  Dtype
---  ---                                ---
0   id                                    59400 non-null  int64
1   amount_tsh                          59400 non-null  float64
2   date_recorded                       59400 non-null  object
3   funder                              55765 non-null  object
4   gps_height                          59400 non-null  int64
5   installer                           55745 non-null  object
6   longitude                           59400 non-null  float64
7   latitude                            59400 non-null  float64
8   wpt_name                            59400 non-null  object
9   num_private                         59400 non-null  int64
10  basin                               59400 non-null  object
11  subvillage                          59029 non-null  object
12  region                              59400 non-null  object
13  region_code                         59400 non-null  int64
14  district_code                      59400 non-null  int64
15  lga                                 59400 non-null  object
16  ward                                59400 non-null  object
17  population                          59400 non-null  int64
18  public_meeting                     56066 non-null  object
19  recorded_by                         59400 non-null  object
20  scheme_management                  55523 non-null  object
21  scheme_name                        31234 non-null  object
22  permit                             56344 non-null  object
23  construction_year                  59400 non-null  int64
24  extraction_type                     59400 non-null  object
25  extraction_type_group               59400 non-null  object
26  extraction_type_class               59400 non-null  object
27  management                         59400 non-null  object
28  management_group                   59400 non-null  object
29  payment                            59400 non-null  object
30  payment_type                       59400 non-null  object
31  water_quality                      59400 non-null  object
32  quality_group                      59400 non-null  object
33  quantity                           59400 non-null  object
34  quantity_group                     59400 non-null  object
35  source                             59400 non-null  object

```

```

36 source_type          59400 non-null object
37 source_class         59400 non-null object
38 waterpoint_type      59400 non-null object
39 waterpoint_type_group 59400 non-null object
40 status_group         59400 non-null object
dtypes: float64(3), int64(7), object(31)
memory usage: 19.0+ MB

```

```
In [8]: merged_data.value_counts()
```

```

Out[8]: id      amount_tsh  date_recorded  funder      gps_height  installer
longitude latitude  wpt_name      num_private  basin      s
ubvillage region      region_code  district_code  lga      ward
population public_meeting recorded_by      scheme_management scheme_n
ame      permit construction_year extraction_type extraction_typ
e_group extraction_type_class management management_group payment
payment_type water_quality quality_group quantity      quantity_group source
source_type source_class waterpoint_type      waterpoint_type_group s
tatus_group
74247  50.0      2013-02-16      Mission      965      DWE
35.432998 -10.639270 Kwa Mapunda      0      Ruvuma / Southern Coast M
pakani      Ruvuma      10      2      Songea Rural Maposeni
900      True      GeoData Consultants Ltd VWC      Mradi wa
maji wa peramiho      True      2009      other      other
other      vwc      user-group      pay per bucket      per
bucket      soft      good      dry      dry      river
river/lake surface      communal standpipe      communal standpipe n
on functional      1
24588  0.0      2013-03-23      Government Of Tanzania 1344      DWE
37.544739 -3.291398 Kwa Bariki Kombe 0      Pangani      B
arazani      Kilimanjaro 3      4      Moshi Rural Mamba Kusini
1      False      GeoData Consultants Ltd VWC      Una mkol
owoni      True      1972      gravity      gravity
gravity      vwc      user-group      never pay      neve
r pay      soft      good      insufficient      insufficient      spring
spring      groundwater      other      other      n
on functional      1
24558  0.0      2011-07-24      Wananchi      0      wananchi
33.814988 -9.490739 Kwa Asukile      0      Lake Nyasa      B
ugoba      Mbeya      12      3      Kyela      Ipande
0      True      GeoData Consultants Ltd VWC      Sinyanga
water supplied sc      True      0      gravity      gravity
gravity      vwc      user-group      never pay      neve
r pay      soft      good      dry      dry      spring
spring      groundwater      communal standpipe      communal standpipe n
on functional      1
24563  0.0      2011-03-14      Go      526      Go
36.990775 -7.400210 Bustanini      0      Rufiji      M
jini      Morogoro      5      1      Kilosa      Mikumi
250      True      GeoData Consultants Ltd Company      Mi
True      1975      gravity      gravity      gravity
company      commercial      never pay      never pay      soft
good      enough      enough      river      river/lake      surface
communal standpipe      communal standpipe      functional      1
24564  0.0      2013-07-03      Government Of Tanzania 1232      RWE
36.874949 -3.343532 Aminieli Nanyaru 0      Pangani      K
iwawa      Arusha      2      7      Meru      Maji ya Chai
120      True      GeoData Consultants Ltd VWC      Tuvaila
gravity water supply True      1968      gravity      gravity
gravity      wug      user-group      unknown      unkn
own      soft      good      enough      enough      river
river/lake surface      communal standpipe      communal standpipe f
unctional      1

```

```
..
```

```

49306 500.0      2013-02-13      Kiwanda Cha Tangawizi  1215      District C
ouncil 37.989865 -4.390224      Kwa Mzee Mkota      0      Pangani
Goha      Kilimanjaro 3      3      Same      Mnyamba
38      True      GeoData Consultants Ltd VWC      Tangawiz
i Water Supply      True      2012      gravity      gravity
gravity      vwc      user-group      pay when scheme fails on f
ailure      soft      good      enough      enough      spring
spring      groundwater      communal standpipe      communal standpipe      f
unctional      1
49308 0.0      2011-04-17      Water      0      DWE
35.915891 -4.695893      Ikova Mwisho      0      Internal      I
kova      Dodoma      1      1      Kondoa      Pahi
0      True      GeoData Consultants Ltd VWC      Pahi
False 0      gravity      gravity      gravity
vwc      user-group      never pay      never pay      soft
good      enough      enough      spring      spring      groundwat
er      communal standpipe      communal standpipe      non functional 1
49310 0.0      2011-04-05      Government Of Tanzania 1488      DWE
38.286117 -4.822377      Kwa Mzee Hoza      0      Pangani      N
yankei      Tanga      4      1      Lushoto      Ubiri
1      True      GeoData Consultants Ltd VWC      Ilente s
treem      True      1981      gravity      gravity
gravity      vwc      user-group      never pay      neve
r pay      soft      good      enough      enough      spring
spring      groundwater      communal standpipe      communal standpipe      n
on functional 1
49311 0.0      2011-03-16      Government Of Tanzania 1793      DWE
34.828328 -9.016824      Kwa Esau Lulambo 0      Rufiji      L
yalamo      Iringa      11      4      Njombe      Mtwango
35      True      GeoData Consultants Ltd VWC      Ilunda p
umping scheme      False 1976      gravity      gravity
gravity      vwc      user-group      never pay      neve
r pay      soft      good      enough      enough      spring
spring      groundwater      communal standpipe      communal standpipe      n
on functional 1
2 0.0      2011-03-27      Lvia      0      LVIA
36.115056 -6.279268      Bombani      0      Wami / Ruvu      S
ongambele      Dodoma      1      4      Chamwino      Msamalo
0      True      GeoData Consultants Ltd VWC      Mgun
True 0      mono      mono      motorpump
vwc      user-group      pay per bucket      per bucket      soft
good      insufficient insufficient machine dbh borehole      groundwat
er      communal standpipe multiple communal standpipe      functional 1
Length: 27813, dtype: int64

```

Clean

Preparing the merged dataset for feature exploration and how they relate to the 'status_group' target variable.

```
In [9]: # Converting 'date_recorded' to datetime
merged_data['date_recorded'] = pd.to_datetime(merged_data['date_recorded'])
```

Addressing Categorical Features with Parent and Subgroup Columns

```
In [10]: grouped = merged_data.groupby(['extraction_type_class', 'extraction_type_group',
print(grouped)
```

extraction_type_class	extraction_type_group	extraction_type	
gravity	gravity	gravity	26780
handpump	afridev	afridev	1770

	india mark ii	india mark ii	2400
	india mark iii	india mark iii	98
	nira/tanira	nira/tanira	8154
	other handpump	other - mkulima/shinyanga	2
		other - play pump	85
		other - swan 81	229
		walimi	48
	swan 80	swan 80	3670
motorpump	mono	mono	2865
	other motorpump	cemo	90
		climax	32
other	other	other	6430
rope pump	rope pump	other - rope pump	451
submersible	submersible	ksb	1415
		submersible	4764
wind-powered	wind-powered	windmill	117
dtype:	int64		

```
In [11]: grouped = merged_data.groupby(['management_group', 'management']).size()
print(grouped)
```

management_group	management	
commercial	company	685
	private operator	1971
	trust	78
	water authority	904
other	other	844
	other - school	99
parastatal	parastatal	1768
unknown	unknown	561
user-group	vwc	40507
	water board	2933
	wua	2535
	wug	6515
dtype:	int64	

```
In [12]: grouped = merged_data.groupby(['waterpoint_type_group', 'waterpoint_type']).size()
print(grouped)
```

waterpoint_type_group	waterpoint_type	
cattle trough	cattle trough	116
communal standpipe	communal standpipe	28522
	communal standpipe multiple	6103
dam	dam	7
hand pump	hand pump	17488
improved spring	improved spring	784
other	other	6380
dtype:	int64	

Handling Missing Values

```
In [13]: # Calculating the percentage of zero values for each column
zero_value_percentages = {}
for column in merged_data.columns:
    zero_count = (merged_data[column] == 0).sum()
    zero_value_percentages[column] = (zero_count / len(merged_data)) * 100

zero_value_percentages
```

```
Out[13]: {'id': 0.0016835016835016834,
'amount_tsh': 70.09932659932659,
'date_recorded': 0.0,
```



```
'funder': 0.0,
'gps_height': 34.40740740740741,
'installer': 0.0,
'longitude': 3.05050505050505,
'latitude': 0.0,
'wpt_name': 0.0,
'num_private': 98.72558922558923,
'basin': 0.0,
'subvillage': 0.0,
'region': 0.0,
'region_code': 0.0,
'district_code': 0.038720538720538725,
'lga': 0.0,
'ward': 0.0,
'population': 35.994949494949495,
'public_meeting': 8.51010101010101,
'recorded_by': 0.0,
'scheme_management': 0.0,
'scheme_name': 0.0,
'permit': 29.44781144781145,
'construction_year': 34.86363636363636,
'extraction_type': 0.0,
'extraction_type_group': 0.0,
'extraction_type_class': 0.0,
'management': 0.0,
'management_group': 0.0,
'payment': 0.0,
'payment_type': 0.0,
'water_quality': 0.0,
'quality_group': 0.0,
'quantity': 0.0,
'quantity_group': 0.0,
'source': 0.0,
'source_type': 0.0,
'source_class': 0.0,
'waterpoint_type': 0.0,
'waterpoint_type_group': 0.0,
'status_group': 0.0}
```

```
In [14]: # Calculating the percentage of missing values in each column
missing_values = merged_data.isnull().mean() * 100
missing_values = missing_values[missing_values > 0].sort_values(ascending=False)

missing_values
```

```
Out[14]: scheme_name      47.417508
scheme_management    6.526936
installer            6.153199
funder               6.119529
public_meeting       5.612795
permit               5.144781
subvillage           0.624579
dtype: float64
```

Since 'scheme_management', 'installer', 'funder', and 'permit' have less than 7% missing values and are potentially relevant, replacing them is a good option.

```
In [15]: # Replacing the missing values
for column in ['scheme_management', 'installer', 'funder', 'permit']:
    merged_data[column].fillna('Unknown', inplace=True)

# Sanity check on missing values
```

```
remaining_missing_values = merged_data.isnull().sum()
remaining_missing_values[remaining_missing_values > 0]
```

```
Out[15]: subvillage      371
public_meeting    3334
scheme_name      28166
dtype: int64
```

Dropping Columns

```
In [16]: # Dropping additional group columns
# Dropping irrelevant columns
# Dropping features with missing values over 45%
features_dropped = merged_data.drop(columns=['quantity_group', 'extraction_type_
                                             'waterpoint_type_group', 'managemen
                                             'payment', 'water_quality', 'source
                                             'region_code', 'public_meeting', 'r
                                             'wpt_name', 'district_code', 'id',
                                             'ward', 'subvillage', 'extraction_t

features_dropped.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 59400 entries, 0 to 59399
Data columns (total 21 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                           59400 non-null  float64
1   date_recorded                        59400 non-null  datetime64[ns]
2   funder                               59400 non-null  object
3   gps_height                           59400 non-null  int64
4   installer                            59400 non-null  object
5   longitude                            59400 non-null  float64
6   latitude                             59400 non-null  float64
7   basin                               59400 non-null  object
8   region                               59400 non-null  object
9   population                           59400 non-null  int64
10  scheme_management                    59400 non-null  object
11  permit                               59400 non-null  object
12  construction_year                    59400 non-null  int64
13  extraction_type_class                 59400 non-null  object
14  management                           59400 non-null  object
15  payment_type                         59400 non-null  object
16  quality_group                        59400 non-null  object
17  quantity                             59400 non-null  object
18  source_type                          59400 non-null  object
19  waterpoint_type                      59400 non-null  object
20  status_group                         59400 non-null  object
dtypes: datetime64[ns](1), float64(3), int64(3), object(14)
memory usage: 10.0+ MB
```

Removing Duplicates

```
In [17]: features_dropped.duplicated().sum()
```

```
Out[17]: 685
```

```
In [18]: data_dedup = features_dropped.drop_duplicates()

# Rechecking for duplicates
```

```
new_duplicate_count = data_dedup.duplicated().sum()
new_duplicate_count
```

Out[18]: 0

Feature Engineering

```
In [19]: # Extracting year and month from 'date_recorded'
data_dedup['year_recorded'] = data_dedup['date_recorded'].dt.year
data_dedup['month_recorded'] = data_dedup['date_recorded'].dt.month

data_dedup[['longitude', 'gps_height', 'construction_year', 'year_recorded', 'month_recorded']]
```

<ipython-input-19-a25e6a562f69>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_dedup['year_recorded'] = data_dedup['date_recorded'].dt.year
```

<ipython-input-19-a25e6a562f69>:3: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_dedup['month_recorded'] = data_dedup['date_recorded'].dt.month
```

Out[19]:

	longitude	gps_height	construction_year	year_recorded	month_recorded
0	34.938093	1390	1999	2011	3
1	34.698766	1399	2010	2013	3
2	37.460664	686	2009	2013	2
3	38.486161	263	1986	2013	1
4	31.130847	0	0	2011	7

```
In [20]: # Adding 'well_age' feature
data_dedup['well_age'] = data_dedup.apply(lambda row: 0 if row['construction_year'] == row['year_recorded']
                                          else row['year_recorded'] - row['construction_year'], axis=1)

# Displaying the first 50 rows to check the 'construction_year' and 'well_age'
data_dedup[['year_recorded', 'construction_year', 'well_age']].value_counts()
```

<ipython-input-20-aef931d43ca2>:2: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_dedup['well_age'] = data_dedup.apply(lambda row: 0 if row['construction_year'] == row['year_recorded']
                                          else row['year_recorded'] - row['construction_year'], axis=1)
```

Out[20]:

year_recorded	construction_year	well_age
2011	0	13104
2012	0	5000
2013	0	1906
	2000	13
	2010	3
		...

2004	2005	-1	1
	2006	-2	1
	2007	-3	1
2012	1984	28	1
2002	0	0	1

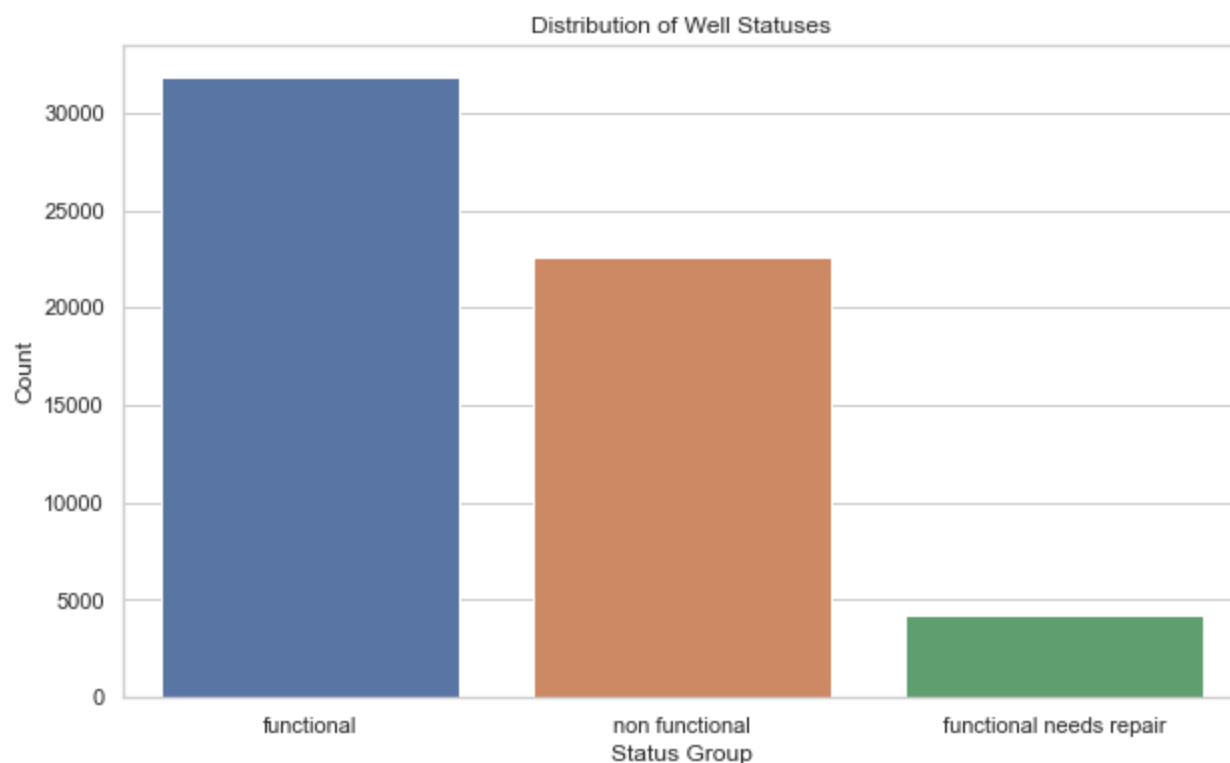
Length: 168, dtype: int64

Creating Binary Target Column

First, we'll need to understand the distribution of the target variable and visualize the proportion of functional vs. non-functional wells.

```
In [21]: sns.set(style="whitegrid")

# Plotting the distribution of well statuses
plt.figure(figsize=(10, 6))
sns.countplot(x='status_group', data=data_dedup)
plt.title('Distribution of Well Statuses')
plt.ylabel('Count')
plt.xlabel('Status Group')
plt.show()
```



To address differing opinions on how to condense our target into a binary column, we will create two separate binary target columns and assess the better performer on our baseline model.

status_binary:

- Class 0 = non-functional & functional needs repair
- Class 1 = functional

status_binary_reversed:

- Class 0 = non-functional

- Class 1 = functional & functional needs repair

```
In [22]: # Binary encoding of the 'status_group' column

# 'functional' is assigned 1 and 'non functional' or 'functional needs repair' a
data_dedup['status_binary'] = data_dedup['status_group'].apply(lambda x: 1 if x

# 'non-functional' is assigned 0 and 'functional' or 'functional needs repair' a
data_dedup['status_binary_reversed'] = data_dedup['status_group'].apply(lambda x

data_dedup[['status_group', 'status_binary', 'status_binary_reversed']].head(20)
```

<ipython-input-22-e016e00ec615>:4: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_dedup['status_binary'] = data_dedup['status_group'].apply(lambda x: 1 if
x == 'functional' else 0)
```

<ipython-input-22-e016e00ec615>:7: SettingWithCopyWarning:
A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
data_dedup['status_binary_reversed'] = data_dedup['status_group'].apply(lambda
x: 0 if x == 'non functional' else 1)
```

```
Out[22]:
```

	status_group	status_binary	status_binary_reversed
--	--------------	---------------	------------------------

0	functional	1	1
1	functional	1	1
2	functional	1	1
3	non functional	0	0
4	functional	1	1
5	functional	1	1
6	non functional	0	0
7	non functional	0	0
8	non functional	0	0
9	functional	1	1
10	functional	1	1
11	functional	1	1
12	functional	1	1
13	functional	1	1
14	functional	1	1
15	functional	1	1
16	non functional	0	0
17	non functional	0	0

	status_group	status_binary	status_binary_reversed
18	functional needs repair	0	1
19	functional	1	1

Creating Master Dataset

```
In [23]: # Creating master dataset with all values of 'well_age' greater than or equal to
# to eliminate negative values where 'recorded_year' was likely listed inaccurat
master_data = data_dedup[data_dedup['well_age'] >= 0]
master_data.info()
```

```
<class 'pandas.core.frame.DataFrame'>
Int64Index: 58706 entries, 0 to 59399
Data columns (total 26 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   amount_tsh                            58706 non-null  float64
1   date_recorded                         58706 non-null  datetime64[ns]
2   funder                                58706 non-null  object
3   gps_height                            58706 non-null  int64
4   installer                             58706 non-null  object
5   longitude                             58706 non-null  float64
6   latitude                              58706 non-null  float64
7   basin                                 58706 non-null  object
8   region                                58706 non-null  object
9   population                            58706 non-null  int64
10  scheme_management                     58706 non-null  object
11  permit                                58706 non-null  object
12  construction_year                     58706 non-null  int64
13  extraction_type_class                  58706 non-null  object
14  management                             58706 non-null  object
15  payment_type                           58706 non-null  object
16  quality_group                          58706 non-null  object
17  quantity                               58706 non-null  object
18  source_type                            58706 non-null  object
19  waterpoint_type                        58706 non-null  object
20  status_group                           58706 non-null  object
21  year_recorded                          58706 non-null  int64
22  month_recorded                         58706 non-null  int64
23  well_age                               58706 non-null  int64
24  status_binary                          58706 non-null  int64
25  status_binary_reversed                 58706 non-null  int64
dtypes: datetime64[ns](1), float64(3), int64(8), object(14)
memory usage: 12.1+ MB
```

```
In [24]: # Saving master dataset to csv
master_data.to_csv('data/master_data.csv', index=False)
```

EDA

Data Distribution Visualizations

Using histograms, we'll plot the distribution of key numeric variables like amount_tsh, population, and gps_height.

```
In [25]: sns.set_style("whitegrid")

# Creating histograms for 'amount_tsh', 'population', and 'gps_height'
```


```
fig, axes = plt.subplots(1, 3, figsize=(18, 6))

# Plot for 'amount_tsh'
sns.histplot(master_data['amount_tsh'], bins=30, ax=axes[0], kde=True)
axes[0].set_title('Distribution of Total Static Head (amount_tsh)')
axes[0].set_xlabel('Amount TSH')
axes[0].set_ylabel('Frequency')

# Plot for 'population'
sns.histplot(master_data['population'], bins=30, ax=axes[1], kde=True)
axes[1].set_title('Distribution of Population')
axes[1].set_xlabel('Population')
axes[1].set_ylabel('Frequency')

# Plot for 'gps_height'
sns.histplot(master_data['gps_height'], bins=30, ax=axes[2], kde=True)
axes[2].set_title('Distribution of GPS Height')
axes[2].set_xlabel('GPS Height')
axes[2].set_ylabel('Frequency')

plt.tight_layout()
plt.show()
```



Observations

These distributions reflect the high number of zero values in our dataset for these features.

Total Static Head (amount_tsh): The distribution appears to be highly skewed to the right, indicating that most wells have a low static head value.

Population: This distribution is also right-skewed, showing that most wells serve a relatively small population, with fewer points serving larger populations.

GPS Height: The distribution is more varied, indicating a range of elevations at which wells are located.

Geographical Analysis

We'll create a geographical plot using latitude and longitude to see if there is any geographical pattern in the status of wells.

In [26]: `!pip install geopandas`

Requirement already satisfied: geopandas in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (0.13.2)

Requirement already satisfied: pyproj>=3.0.1 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (3.5.0)

Requirement already satisfied: fiona>=1.8.19 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (1.9.5)

Requirement already satisfied: pandas>=1.1.0 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (1.1.3)

Requirement already satisfied: shapely>=1.7.1 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (2.0.2)

Requirement already satisfied: packaging in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from geopandas) (20.4)

Requirement already satisfied: certifi in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pyproj>=3.0.1->geopandas) (2023.7.22)

Requirement already satisfied: setuptools in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (50.3.0.post20201103)

Requirement already satisfied: six in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (1.15.0)

Requirement already satisfied: click~8.0 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (8.1.7)

Requirement already satisfied: click-plugins>=1.0 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (1.1.1)

Requirement already satisfied: cligj>=0.5 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (0.7.2)

Requirement already satisfied: importlib-metadata; python_version < "3.10" in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (2.0.0)

Requirement already satisfied: attrs>=19.2.0 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from fiona>=1.8.19->geopandas) (20.2.0)

Requirement already satisfied: python-dateutil>=2.7.3 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=1.1.0->geopandas) (2.8.1)

Requirement already satisfied: pytz>=2017.2 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=1.1.0->geopandas) (2020.1)

Requirement already satisfied: numpy>=1.15.4 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from pandas>=1.1.0->geopandas) (1.18.5)

Requirement already satisfied: pyparsing>=2.0.2 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from packaging->geopandas) (2.4.7)

Requirement already satisfied: zipp>=0.5 in /Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages (from importlib-metadata; python_version < "3.10"->fiona>=1.8.19->geopandas) (3.3.0)

```
In [30]: import geopandas
import matplotlib.pyplot as plt

# Create GeoDataFrame
gdf = geopandas.GeoDataFrame(
    master_data, geometry=geopandas.points_from_xy(master_data.longitude, master_data.latitude)

# Assigning 'status_group'
functional = gdf[gdf['status_group'] == 'functional']
repair = gdf[gdf['status_group'] == 'functional needs repair']
non_functional = gdf[gdf['status_group'] == 'non functional']

# Load world shapefile
world_shapefile_path = 'data/ne_110m_admin_0_countries/ne_110m_admin_0_countries.shp'
world = geopandas.read_file(world_shapefile_path)

# Filter for Tanzania
```



```

fig, ax = plt.subplots(figsize=(15, 15))
base = world[world.ADMIN == 'United Republic of Tanzania'].plot(color='white', e

# Scatter plots for each category
ax.scatter(functional['longitude'], functional['latitude'], c='lightblue', alpha=0.5)
ax.scatter(repair['longitude'], repair['latitude'], c='orange', alpha=1, s=3, label='Needs Repair')
ax.scatter(non_functional['longitude'], non_functional['latitude'], c='darkblue', alpha=1, s=3, label='Non Functional')

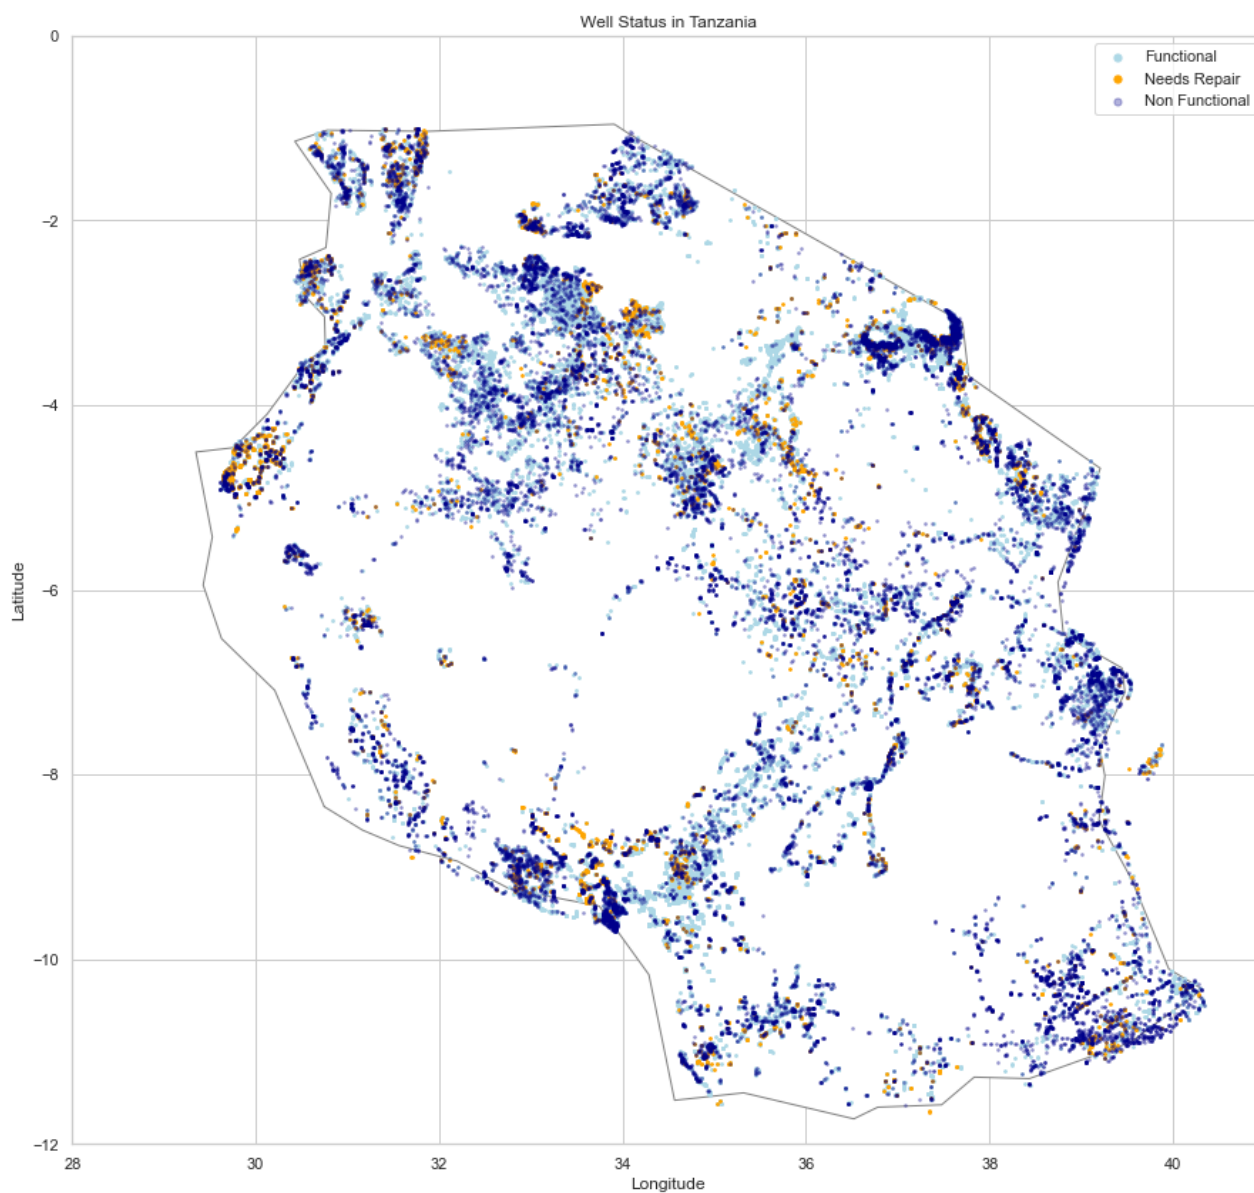
# Limiting the display area
ax.set_ylim(-12, 0)
ax.set_xlim(28, 41)

# Adding labels and title
ax.set_xlabel('Longitude')
ax.set_ylabel('Latitude')
ax.set_title('Well Status in Tanzania')

# Adding legend
ax.legend(markerscale=3, loc='upper right')

plt.show()

```



To further identify geographic feature relationships, we'll plot only the non-functional wells and adjust the marker size by 'gps_height'.

```
In [32]: # Filter for 'non functional' wells
non_functional = gdf[gdf['status_group'] == 'non functional']

# Filter for Tanzania and plot
fig, ax = plt.subplots(figsize=(15, 15))
base = world[world.ADM1 == 'United Republic of Tanzania'].plot(color='white', e

# Plot for 'non functional' wells with marker size based on 'gps_height'
# Normalize 'gps_height' for visualization
max_height = non_functional['gps_height'].max()
marker_size = (non_functional['gps_height'] / max_height) * 100

ax.scatter(non_functional['longitude'], non_functional['latitude'], c='blue', al


# Limiting the display area
ax.set_ylim(-12, 0)
ax.set_xlim(28, 41)

# Adding labels and title
ax.set_xlabel('Longitude')
ax.set_ylabel('Latitude')
ax.set_title('Non Functional Wells in Tanzania')

# Adding legend
ax.legend(markerscale=3, loc='upper right')

plt.show()
```

```
/Users/kariprimiano/anaconda3/envs/learn-env/lib/python3.8/site-packages/matplot
lib/collections.py:922: RuntimeWarning: invalid value encountered in sqrt
scale = np.sqrt(self._sizes) * dpi / 72.0 * self._factor
```



The majority of 'non functional' wells seem to be located at higher altitudes.

Categorical to Target Relationships

We'll explore the relationships between categorical variables and the target variable 'status_group'.

```
In [36]: import math

categorical_vars = ['basin', 'region', 'source_type', 'quality_group', 'extracti

# Define custom colors for each status
color_map = {'functional': 'blue', 'non functional': 'darkblue', 'functional nee

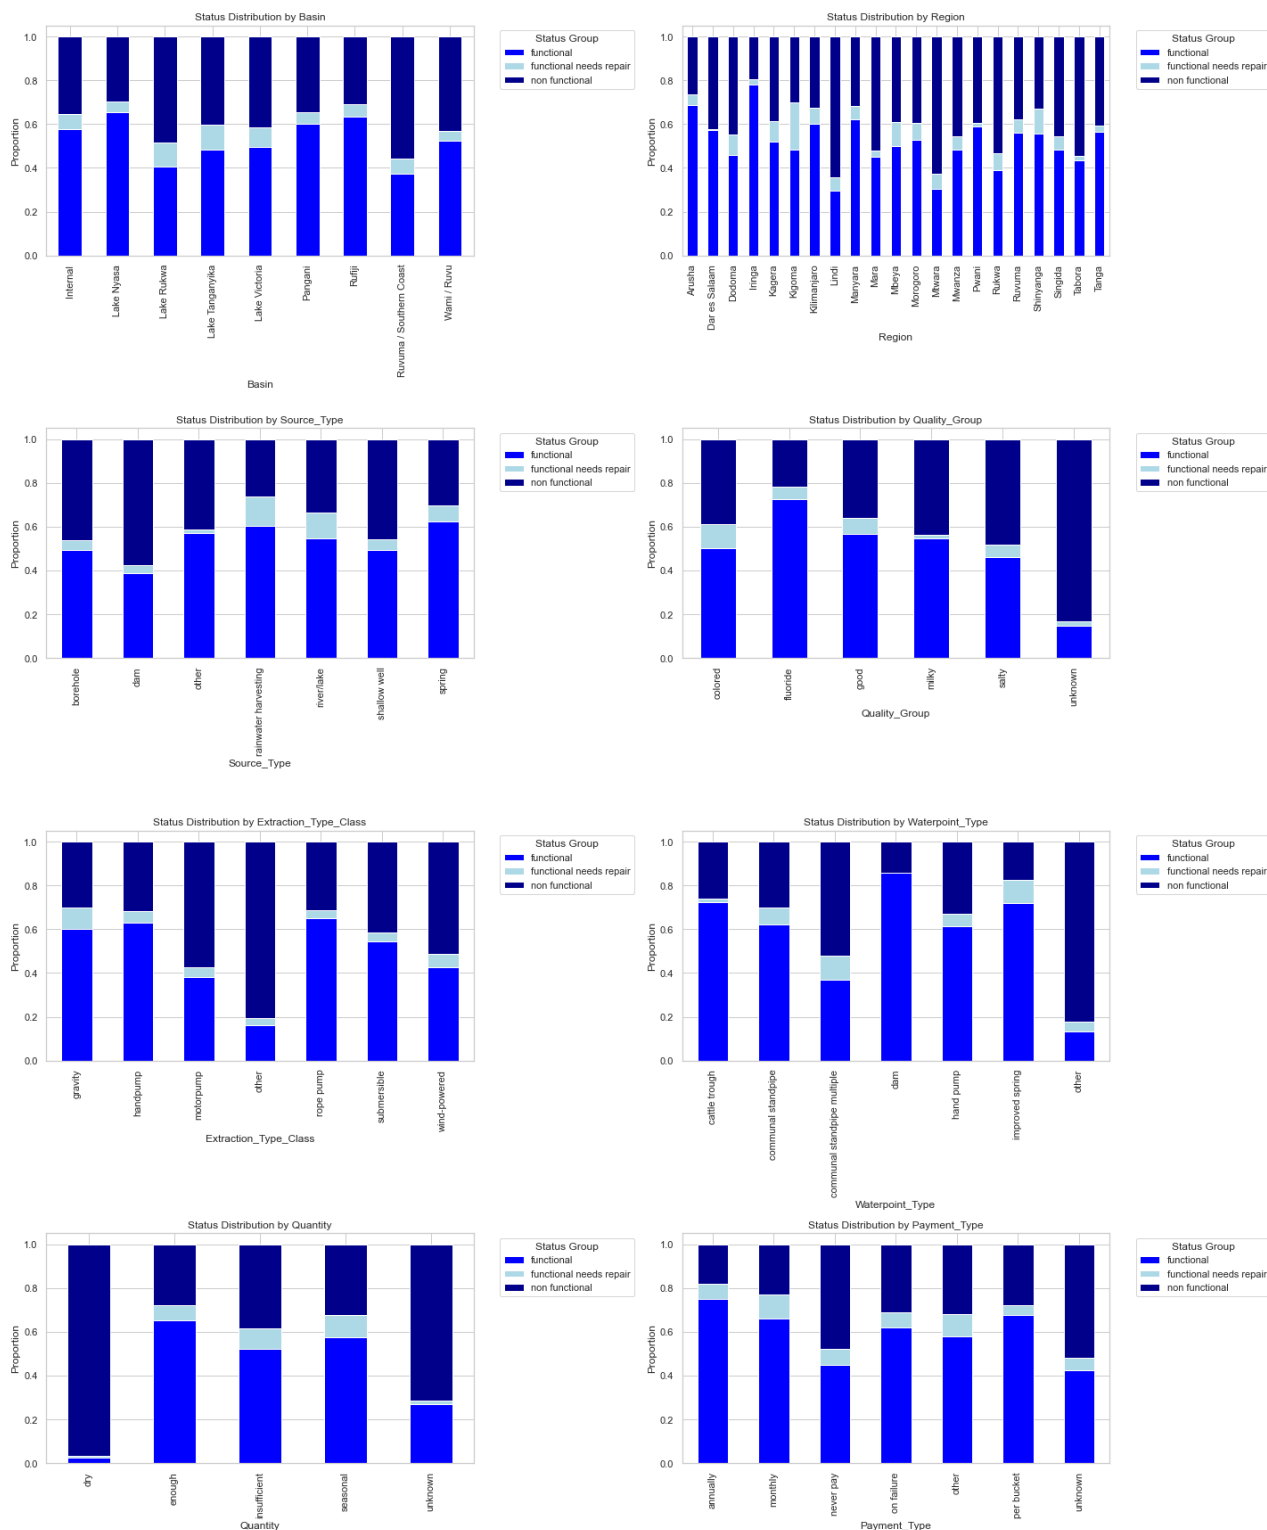
# Determine the number of rows needed for the subplot (2 plots per row)
num_vars = len(categorical_vars)
num_rows = math.ceil(num_vars / 2)

# Creating subplots for each categorical variable
fig, axes = plt.subplots(nrows=num_rows, ncols=2, figsize=(20, num_rows * 6))
axes = axes.flatten() # Flatten the axes array for easy iteration
```

```
# Looping through the variables and creating a stacked bar plot for each
for i, var in enumerate(categorical_vars):
    # Creating a crosstab for the variable and status_group
    crosstab = pd.crosstab(master_data[var], master_data['status_group'], normal

    # Creating a stacked bar plot with custom colors
    crosstab.plot(kind='bar', stacked=True, color=[color_map[status] for status
    axes[i].set_title(f'Status Distribution by {var.title()}')
    axes[i].set_xlabel(var.title())
    axes[i].set_ylabel('Proportion')
    axes[i].legend(title='Status Group', bbox_to_anchor=(1.05, 1), loc='upper le

# Adjust the layout
plt.tight_layout()
plt.show()
```



Observations

Region: Similar to basins, each region has a unique distribution of well statuses. This seems to be an indicator for well status.

Payment Type: Whether a well is paid seems to be a crucial factor. Wells that are not paid have a high number of non-functional wells.

Waterpoint Type: The method used for the population to access the water from the wells is another crucial factor. Similarly to extraction methods, waterpoint types might be more robust

and less prone to failure, while others could be more complex and require frequent repairs.

Machine Learning

Baseline Model #1: Binary Target Column

Class 0 = non-functional/needs repair

Class 1 = functional

In the context of non-functional wells, focusing on recall (false negative) may be more important to ensure that most of the non-functional wells are correctly identified. We will test two separate baseline models, each with a different binary target column to see which performs best on recall.

```
In [37]: # Initiating train_test_split
X = master_data[['waterpoint_type']]
y = master_data['status_binary']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random

# Encoding categorical variable
from sklearn.preprocessing import OneHotEncoder
ohe = OneHotEncoder()
ohe.fit(X_train)
X_train_encoded = ohe.transform(X_train)
X_test_encoded = ohe.transform(X_test)

# Plotting Log Reg transform
logreg = LogisticRegression(random_state=42)
logreg.fit(X_train_encoded, y_train)

# Checking if the target is balanced
y_test.value_counts(normalize=True)

y_pred = logreg.predict(X_test_encoded)

print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.75         0.35         0.47         8894
     1       0.62         0.90         0.73        10479

 accuracy          0.69
 macro avg         0.69         0.62         0.60         19373
weighted avg         0.68         0.65         0.61         19373
```

```
Confusion Matrix:
[[3070 5824]
 [1015 9464]]
```

Evaluating with ROC Curve

```
In [38]: # Predict probas for the positive class
y_pred_proba = logreg.predict_proba(X_test_encoded)[: , 1]
```


```

# Compute AUC-ROC
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

# Plotting ROC Curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

```



Baseline Model #2: Reverse Binary Target Column

Class 0 = non-functional

Class 1 = functional/needs repair

```

In [39]: # Intitating train_test_split
X = master_data[['waterpoint_type']]
y = master_data['status_binary_reversed']
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.33, random

# Encoding categorical variable
from sklearn.preprocessing import OneHotEncoder
ohe = OneHotEncoder()
ohe.fit(X_train)
X_train_encoded = ohe.transform(X_train)
X_test_encoded = ohe.transform(X_test)

# Plotting Log Reg transform
logreg = LogisticRegression(random_state=42)
logreg.fit(X_train_encoded, y_train)

# Checking if the target is balanced
y_test.value_counts(normalize=True)

```

```
# Predicting and evaluating the model
y_pred = logreg.predict(X_test_encoded)
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.68       0.37       0.48       7490
     1       0.69       0.89       0.78      11883

 accuracy          0.69       19373
 macro avg       0.68       0.63       0.63       19373
 weighted avg    0.68       0.69       0.66       19373
```

```
Confusion Matrix:
[[ 2758  4732]
 [ 1327 10556]]
```


Evaluating with ROC Curve

```
In [40]: # Predict probas for the positive class
y_pred_proba = logreg.predict_proba(X_test_encoded)[:, 1]

# Compute AUC-ROC
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

# Plotting ROC Curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



Baseline Comparison

Baseline 1

Precision:

- Class 0: 0.75 (Higher precision for class 0)
- Class 1: 0.62

Recall:

- Class 0: 0.35 (Lower recall for class 0)
- Class 1: 0.90 (Higher recall for class 1)

F1-Score:

- Class 0: 0.47 (Lower F1-score for class 0)
- Class 1: 0.73

Accuracy: 65%

Baseline 2

Precision:

- Class 0: 0.68 (Lower precision for class 0)
- Class 1: 0.69

Recall:

- Class 0: 0.37 (Slightly higher recall for class 0)
- Class 1: 0.89 (Slightly lower recall for class 1)

F1-Score:

- Class 0: 0.48 (Slightly higher F1-score for class 0)
- Class 1: 0.78

Accuracy: 69% (Higher)

Analysis

Baseline 2 shows improved overall performance, with better accuracy and a better balance in precision and recall for both classes. However, it is more prone to falsely identifying class 0 (non-functional) instances as class 1 (functional).

Baseline 1 while having higher precision for class 0 (non-functional), falls short in accurately identifying class 0 (non-functional) instances (lower recall).

Since we are more concerned with better recall, we will continue our modeling with **Baseline 2**.

Random Forest Classifier

This model can provide insights into the importance of various features in predicting well functionality. It's less likely to overfit than individual decision trees and doesn't require feature scaling.

```
In [45]: X = master_data.drop(['status_binary', 'status_binary_reversed', 'status_group',
                             'installer', 'permit', 'date_recorded', 'construction_year'])
y = master_data['status_binary_reversed']

# Initializing train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

# Identifying categoricals
categorical_cols = X_train.select_dtypes(include=['object', 'category']).columns

# Creating a column transformer with OneHotEncoder for categoricals
column_transformer = ColumnTransformer(
    transformers=[
        ('cat', OneHotEncoder(handle_unknown='ignore'), categorical_cols)
    ],
    remainder='passthrough'
)

# Applying the column transformer
X_train_encoded = column_transformer.fit_transform(X_train)
X_test_encoded = column_transformer.transform(X_test)

# Creating and training the Random Forest model
rf_model = RandomForestClassifier(n_estimators=100, random_state=42)
rf_model.fit(X_train_encoded, y_train)

# Extracting feature names manually for older versions of scikit-learn
onehot_features = column_transformer.named_transformers_['cat'].get_feature_name
other_features = [col for col in X_train.columns if col not in categorical_cols]
feature_names = np.concatenate([onehot_features, other_features])

# Predicting and evaluating the model
y_pred = rf_model.predict(X_test_encoded)
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.83         0.77         0.80         6856
     1       0.86         0.90         0.88        10756

 accuracy          0.85
 macro avg         0.85
 weighted avg      0.85
```

```
Confusion Matrix:
[[5290 1566]
 [1087 9669]]
```


Evaluating with ROC Curve

```
In [42]: # Predict probas for the positive class
y_pred_proba = rf_model.predict_proba(X_test_encoded)[: , 1]

# Compute AUC-ROC
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

# Plotting ROC Curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



Evaluating Feature Importance

```
In [43]: # Extracting and displaying feature importances
importances = rf_model.feature_importances_
importance_df = pd.DataFrame({'Feature': feature_names, 'Importance': importance_
top_features = importance_df.sort_values(by='Importance', ascending=False).head(
print(top_features)
```

	Feature	Importance
84	latitude	0.153992
83	longitude	0.152580
62	quantity_dry	0.101336
82	gps_height	0.077973
85	population	0.054906
80	waterpoint_type_other	0.040619
81	amount_tsh	0.030302
33	extraction_type_class_other	0.030182
63	quantity_enough	0.028917
87	month_recorded	0.026894
64	quantity_insufficient	0.015844
51	payment_type_never pay	0.013426

75	waterpoint_type_communal_standpipe	0.011889
30	extraction_type_class_gravity	0.010953
44	management_vwc	0.010815

Tuning Random Forest Classifier

- SMOTE for oversampling the minority class or adjusting class weights in the model.
- Hyperparameter tuning of the Random Forest model
- RandomizedSearchCV will randomly sample 10 combos of parameters and use 3-fold cross-validation. This will reduce run time compared to GridSearchCV

```
In [44]: from imblearn.over_sampling import SMOTE
from sklearn.model_selection import RandomizedSearchCV

# Handling class imbalance with SMOTE
smote = SMOTE()
X_train_resampled, y_train_resampled = smote.fit_resample(X_train_encoded, y_train_encoded)

# Define the hyperparameter grid
param_grid = {
    'n_estimators': [100, 200, 300],
    'max_depth': [10, 20, 30],
    'min_samples_split': [2, 5, 10],
    'min_samples_leaf': [1, 2, 4],
    'max_features': ['sqrt', 'log2']
}

# Hyperparameter tuning with Randomized Search
random_search = RandomizedSearchCV(
    RandomForestClassifier(random_state=42),
    param_grid,
    n_iter=10,
    cv=3,
    scoring='recall',
    n_jobs=-1
)

random_search.fit(X_train_resampled, y_train_resampled)

# Get the best model
best_model = random_search.best_estimator_

# Re-train and evaluate the model with the best params
best_model.fit(X_train_resampled, y_train_resampled)
y_pred = best_model.predict(X_test_encoded)

# Evaluate the model
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))
```

```
Classification Report:
              precision    recall  f1-score   support

     0       0.84         0.62         0.71         6856
     1       0.79         0.93         0.85        10756

 accuracy                   0.81         17612
```

macro avg	0.82	0.77	0.78	17612
weighted avg	0.81	0.81	0.80	17612

Confusion Matrix:

```
[[4227 2629]
 [ 787 9969]]
```


Evaluating with ROC Curve

```
In [46]: # Predict probas for the positive class
y_pred_proba = rf_model.predict_proba(X_test_encoded)[: , 1]

# Compute AUC-ROC
roc_auc = roc_auc_score(y_test, y_pred_proba)

# Compute ROC curve
fpr, tpr, _ = roc_curve(y_test, y_pred_proba)

# Plotting ROC Curve
plt.figure()
plt.plot(fpr, tpr, color='darkorange', lw=2, label='ROC curve (area = %0.2f)' %
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()
```



Comparison Interpretations of Random Forest Models

Class 0 = non-functional/needs repair

Class 1 = functional

Precision:

- **Original Model:** Precision is 0.83 for both classes.
- **Tuned Model:** Precision increased to 0.85 for class 0 but decreased to 0.79 for class 1.

Recall:

- **Original Model:** Recall is 0.77 for class 0 and 0.90 for class 1.
- **Tuned Model:** Recall decreased to 0.62 for class 0 but increased to 0.93 for class 1.

F1-Score:

- **Original Model:** F1-scores are 0.80 (class 0) and 0.88 (class 1).
- **Tuned Model:** F1-scores are 0.72 (class 0) and 0.86 (class 1).

Accuracy:

- **Original Model:** Overall accuracy is 0.85.
- **Tuned Model:** Overall accuracy decreased to 0.81.

Macro and Weighted Averages:

- **Original Model:** Both macro and weighted averages are around 0.85.
- **Tuned Model:** Both macro and weighted averages are around 0.77 - 0.81.

AUC-ROC Score:

- **Original Model:** The area under the curve is 0.91, which is high. This means the model can effectively distinguish between the positive class (class 1) and the negative class (class 0).
- **Tuned Model:** The area under the curve is also 0.91.

Analysis of Comparison:

- Tuning the model appears to have made it more biased towards class 1, improving its ability to detect class 1 instances but worsening its performance for class 0 (higher false positives).
- The original model is more balanced in terms of precision and recall across both classes.
- The tuned model has a lower overall accuracy compared to the original model.

Logistic Regression Model: Most Important Features

```
In [47]: # Selecting most important features
features = ['region', 'quantity', 'gps_height', 'waterpoint_type', 'extraction_t
X = master_data[features]
y = master_data['status_binary_reversed']

# Initializing train/test split
X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.3, random_

# Identifying categorical and numerical
```

```

categorical_cols = ['quantity', 'region', 'waterpoint_type', 'extraction_type_cl
numerical_cols = ['gps_height']

# Creating a column transformer with OneHotEncoder and StandardScaler
preprocessor = ColumnTransformer(
    transformers=[
        ('num', StandardScaler(), numerical_cols),
        ('cat', OneHotEncoder(), categorical_cols)
    ]
)

# Creating a pipeline with preprocessing and log reg model
model_pipeline = Pipeline([
    ('preprocessor', preprocessor),
    ('classifier', LogisticRegression(solver='saga', max_iter=1000, random_state
])

# Training the model
model_pipeline.fit(X_train, y_train)

# Predicting and evaluating the model
y_pred = model_pipeline.predict(X_test)
print("Classification Report:\n", classification_report(y_test, y_pred))
print("Confusion Matrix:\n", confusion_matrix(y_test, y_pred))

```

Classification Report:

	precision	recall	f1-score	support
0	0.83	0.54	0.65	6856
1	0.76	0.93	0.84	10756
accuracy			0.78	17612
macro avg	0.79	0.73	0.74	17612
weighted avg	0.79	0.78	0.76	17612

Confusion Matrix:

```

[[3694 3162]
 [ 768 9988]]


```

```

In [48]: # Calculate probas, ROC curve, and AUC for log reg
logreg_probas = model_pipeline.predict_proba(X_test)[: , 1]
fpr_logreg, tpr_logreg, _ = roc_curve(y_test, logreg_probas)
roc_auc_logreg = auc(fpr_logreg, tpr_logreg)

# Plotting ROC Curve
plt.figure()
plt.plot(fpr_logreg, tpr_logreg, color='darkorange', lw=2, label='Logistic Regre
plt.plot([0, 1], [0, 1], color='navy', lw=2, linestyle='--')
plt.xlim([0.0, 1.0])
plt.ylim([0.0, 1.05])
plt.xlabel('False Positive Rate')
plt.ylabel('True Positive Rate')
plt.title('Receiver Operating Characteristic')
plt.legend(loc="lower right")
plt.show()

```



Interpretation

Precision:

- For class 0: 83% precision means that out of all instances predicted as class 0, 83% were actually class 0.
- For class 1: 76% precision indicates that out of all instances predicted as class 1, 76% were actually class 1.

Recall:

- For class 0: The recall of 54% is moderate, meaning the model correctly identifies 54% of the actual class 0 instances.
- For class 1: A high recall of 93% indicates the model is very effective at identifying class 1 instances.

F1-Score:

- For class 0: The F1-score of 0.65 suggests a balance between precision and recall for class 0, but more weighted towards precision.
- For class 1: The F1-score of 0.84 indicates a strong balance between precision and recall for class 1, favoring recall.

Accuracy:

- The overall accuracy of 78% indicates that the model correctly predicts the class for 78% of all instances.

Macro and Weighted Averages:

- Macro average treats both classes equally, showing an average precision of 79%, recall of 73%, and F1-score of 74%.

- Weighted average considers class imbalance, showing slightly higher precision and recall, indicating better performance on the more prevalent class 1.

ROC Score:

- An ROC score of 0.82 suggests a good ability of the model to distinguish between the two classes. It indicates a favorable balance between the true positive rate and false positive rate across different thresholds.

Insights:

- The model performs well overall, especially in predicting class 1, which is indicated by the high recall and F1-score for class 1.
- Model is less effective in correctly identifying class 0 instances, as evidenced by the lower recall for class 0.
- The relatively high number of false positives for class 0 (3162) indicates that the model often misclassifies class 1 instances as class 0.

Key Findings

- **Geographic Indicators:** Including region and altitude, geographic features are 21% MORE influential in identifying non-functional wells than other features.
- **Region:** Mbeya, Morogoro, and Kilimanjaro have the highest rates of non-functional wells. Altitude may play an important role in water source access.
- **Type of Wells:** Communal Standpipe wells are most likely to be a functional well. Other well types have the highest percentage of non-functional wells at 81.38%. Multi Communal Handpipe wells have the second highest percentage of non-functional wells at 53.85%.
- **Payment Type:** Whether a well is paid seems to be a crucial factor. Wells that are not paid have a high number of non-functional wells.
- **Random Forest Classifier:** Our best performing model gave us actionable insights into feature importance and effectively minimized false-negatives.

Conclusion

In this project, we have unearthed critical insights to steer the strategic decision-making of the Tanzanian Government. This includes pinpointing the precise types of wells that warrant prioritized construction efforts, as well as identifying the specific regions that should receive initial focus and substantial investment in well infrastructure.

Next Steps

- **Investigate Additional Features:** Concentrating on geographical indicators like climate, population, and amount of water available in the area.
- **Time-Series Analysis:** Further consideration of the well ages should be analyzed to predict the average lifetime of more robust well structures.
- **Repairs:** Local governments should look at what type of water wells are needing repairs, and the severity of those repairs, to fine-tune non-functional indicators.

Sources

- [Driven Data - Tanzanian Water Wells](#)
 - [Labels](#)
 - [Values](#)
- [The World Bank](#)
- [Groundwater Wells](#)
- [Detection of Non-Function Bore Wells Using Machine Learning Algorithms](#)

About Us

Goknur Kaya

- Github Lead
- Email: goknurkaya@gmail.com
- github.com/goknurk

Kari Primiano

- Tech Lead
- Email: kkprim@gmail.com
- github.com/kkprim

Mytreyi Abburu

- Presentation Lead
- Email: abburumk@gmail.com
- github.com/myt-hue